

```
In [1]: import pandas as pd
import numpy as np
import matplotlib as mpl
import matplotlib.pyplot as plt
import seaborn as sns
mpl.style.use(['ggplot'])
%matplotlib inline
```

```
In [2]: from sklearn import preprocessing, metrics
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import PolynomialFeatures, StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import r2_score, mean_squared_error
from sklearn.neighbors import KNeighborsClassifier
```

```
In [3]: Wdata = pd.read_csv('seattle-weather.csv')
```

```
In [4]: Wdata.head()
```

```
Out[4]:
```

	date	precipitation	temp_max	temp_min	wind	weather
0	2012-01-01	0.0	12.8	5.0	4.7	drizzle
1	2012-01-02	10.9	10.6	2.8	4.5	rain
2	2012-01-03	0.8	11.7	7.2	2.3	rain
3	2012-01-04	20.3	12.2	5.6	4.7	rain
4	2012-01-05	1.3	8.9	2.8	6.1	rain

```
In [5]: Wdata.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1461 entries, 0 to 1460
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   date            1461 non-null  object
1   precipitation    1461 non-null  float64
2   temp_max        1461 non-null  float64
3   temp_min        1461 non-null  float64
4   wind            1461 non-null  float64
5   weather         1461 non-null  object
dtypes: float64(4), object(2)
memory usage: 68.6+ KB
```

```
In [6]: Wdata.date=pd.to_datetime(Wdata.date)
```

```
In [7]: Wdata.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1461 entries, 0 to 1460
Data columns (total 6 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   date            1461 non-null   datetime64[ns]
 1   precipitation    1461 non-null   float64
 2   temp_max        1461 non-null   float64
 3   temp_min        1461 non-null   float64
 4   wind            1461 non-null   float64
 5   weather         1461 non-null   object
dtypes: datetime64[ns](1), float64(4), object(1)
memory usage: 68.6+ KB
```

```
In [8]: Wdata.shape
```

```
Out[8]: (1461, 6)
```

```
In [11]: Wdata.columns
```

```
Out[11]: Index(['date', 'precipitation', 'temp_max', 'temp_min', 'wind', 'weather'], dtype='object')
```

```
In [12]: Wdata.isna().sum()
```

```
Out[12]: date            0
precipitation          0
temp_max               0
temp_min               0
wind                   0
weather                0
dtype: int64
```

```
In [15]: Wdata.duplicated().sum()
```

```
Out[15]: 0
```


```
In [16]: Wdata['weather'].value_counts()
```

```
Out[16]: rain           641
sun             640
fog             101
drizzle         53
```

snow 26
Name: weather, dtype: int64


```
In [19]: wdata['weather'].value_counts().plot(kind='pie',  
figsize=(14,8),  
autopct='%1.1f%%',  
label=True)
```

Out[19]: <AxesSubplot:ylabel='True'>




```
In [24]: wdata['temp_max'].value_counts().plot(kind='density', figsize=(14,8))
```

Out[24]: <AxesSubplot:ylabel='Density'>




```
In [56]: wdata.groupby(['date'])['temp_max'].max().nlargest(10).plot(kind='bar', figsize=(10,8))
```

```
Out[56]: <AxesSubplot: xlabel='date'>
```



```
In [58]: import warnings
warnings.filterwarnings('ignore')
sns.countplot("weather", data=wdata, palette="hls")
```

```
Out[58]: <AxesSubplot:xlabel='weather', ylabel='count'>
```



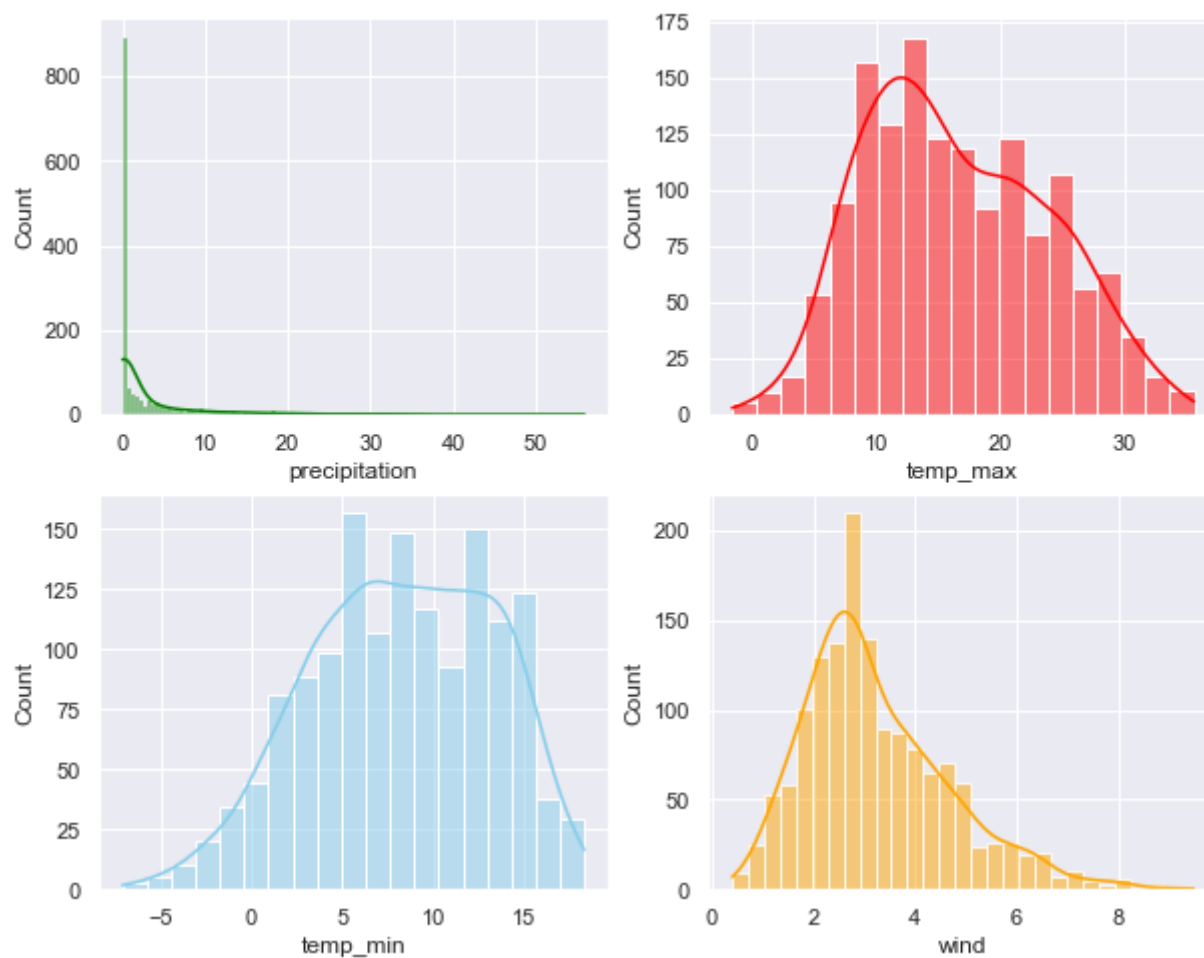
```
In [60]: Wdata.describe()
```

```
Out[60]:
```

	precipitation	temp_max	temp_min	wind
count	1461.000000	1461.000000	1461.000000	1461.000000
mean	3.029432	16.439083	8.234771	3.241136
std	6.680194	7.349758	5.023004	1.437825
min	0.000000	-1.600000	-7.100000	0.400000
25%	0.000000	10.600000	4.400000	2.200000
50%	0.000000	15.600000	8.300000	3.000000
75%	2.800000	22.200000	12.200000	4.000000
max	55.900000	35.600000	18.300000	9.500000

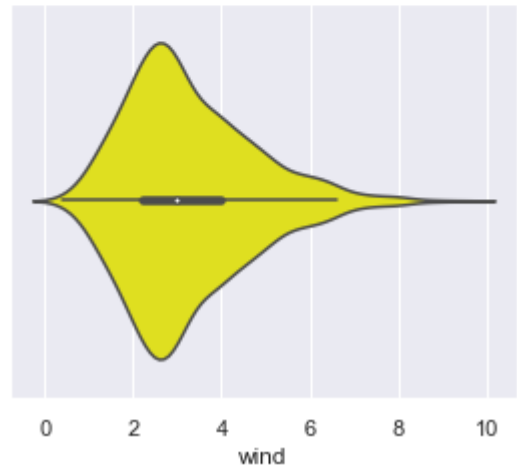



```
In [62]: sns.set(style="darkgrid")
fig,axs=plt.subplots(2,2,figsize=(10,8))
sns.histplot(data=Wdata,x="precipitation",kde=True,ax=axs[0,0],color='green')
sns.histplot(data=Wdata,x="temp_max",kde=True,ax=axs[0,1],color='red')
sns.histplot(data=Wdata,x="temp_min",kde=True,ax=axs[1,0],color='skyblue')
sns.histplot(data=Wdata,x="wind",kde=True,ax=axs[1,1],color='orange')
```

```
Out[62]: <AxesSubplot:xlabel='wind', ylabel='Count'>
```




```
In [64]: sns.set(style="darkgrid")
fig, axs=plt.subplots(2,2,figsize=(10,8))
sns.violinplot(data=Wdata,x="precipitation",kde=True,ax=axs[0,0],color='green')
sns.violinplot(data=Wdata,x="temp_max",kde=True,ax=axs[0,1],color='red')
sns.violinplot(data=Wdata,x="temp_min",kde=True,ax=axs[1,0],color='skyblue')
sns.violinplot(data=Wdata,x="wind",kde=True,ax=axs[1,1],color='yellow')
```

```
Out[64]: <AxesSubplot: xlabel='wind'>
```




```
In [66]: plt.figure(figsize=(12,6))
sns.boxplot("precipitation", "weather", data=Wdata, palette="YlOrBr")
```

```
Out[66]: <AxesSubplot:xlabel='precipitation', ylabel='weather'>
```


```
In [68]: plt.figure(figsize=(12,6))  
sns.boxplot("temp_max", "weather", data=Wdata, palette="inferno")
```

```
Out[68]: <AxesSubplot:xlabel='temp_max', ylabel='weather'>
```




```
In [69]: plt.figure(figsize=(12,6))  
sns.boxplot("wind", "weather", data=wdata, palette="inferno")
```

```
Out[69]: <AxesSubplot:xlabel='wind', ylabel='weather'>
```



```
In [70]: plt.figure(figsize=(12,7))  
sns.heatmap(wdata.corr(), annot=True, cmap='coolwarm')
```

```
Out[70]: <AxesSubplot:>
```




```
In [74]: from scipy import stats
```

```
In [75]: Wdata.plot("precipitation", "temp_max", style='o')
print("Pearson correlation:", Wdata["precipitation"].corr(Wdata["temp_max"]))
print("T Test and P value:", stats.ttest_ind(Wdata["precipitation"], Wdata["temp_max"]))
```


Pearson correlation: -0.22855481643297043

T Test and P value: Ttest_indResult(statistic=-51.60685279531918, pvalue=0.0)



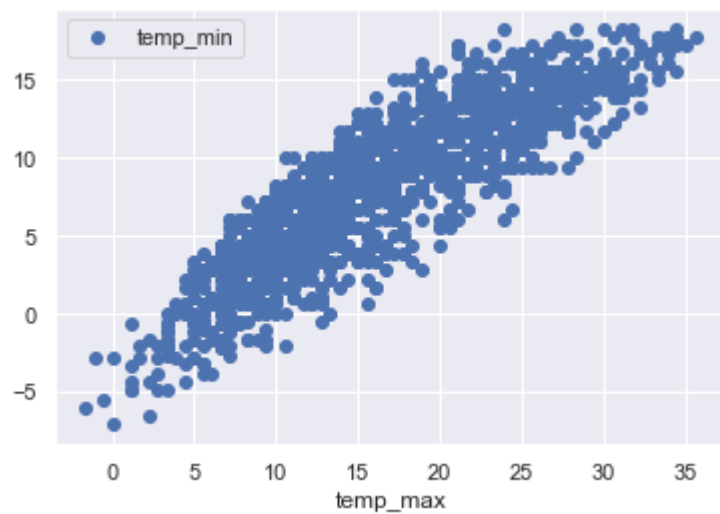
```
In [78]: Wdata.plot("wind", "temp_max", style='o')
print("Pearson correlation:", Wdata["wind"].corr(Wdata["temp_max"]))
print("T Test and P value:", stats.ttest_ind(Wdata["wind"], Wdata["temp_max"]))
```

Pearson correlation: -0.1648566348749548
T Test and P value: Ttest_indResult(statistic=-67.3601643301846, pvalue=0.0)



```
In [79]: Wdata.plot("temp_max", "temp_min", style='o')
```

```
Out[79]: <AxesSubplot:xlabel='temp_max'>
```



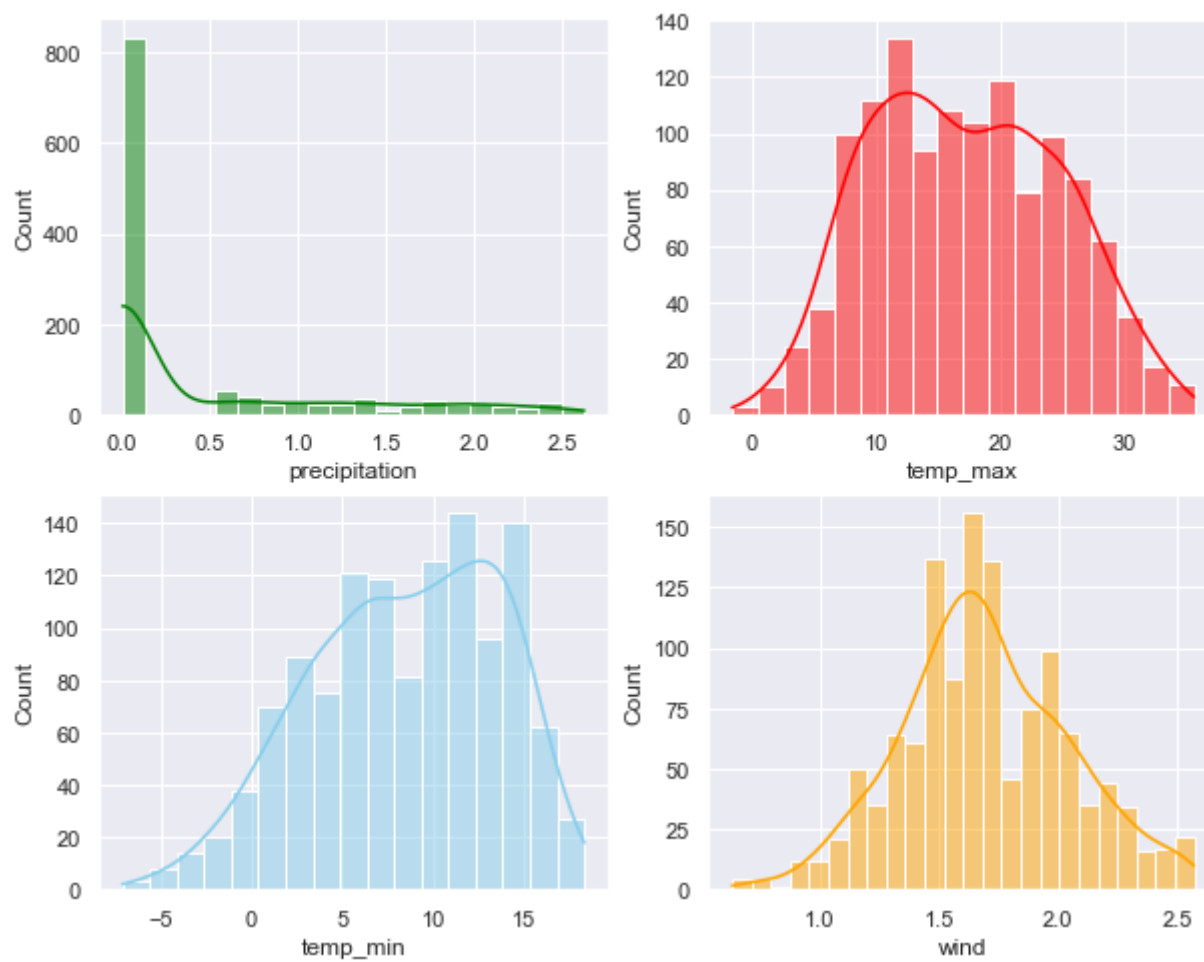
```
In [82]: df=Wdata.drop(["date"], axis=1)
```

```
In [83]: Q1=df.quantile(0.25)
Q3=df.quantile(0.75)
IQR=Q3-Q1
df=df[~((df<(Q1-1.5*IQR))|(df>(Q3+1.5*IQR))).any(axis=1)]
```

```
In [84]: df.precipitation=np.sqrt(df.precipitation)
df.wind=np.sqrt(df.wind)
```

```
In [85]: sns.set(style="darkgrid")
fig,axs=plt.subplots(2,2,figsize=(10,8))
sns.histplot(data=df,x="precipitation",kde=True,ax=axs[0,0],color='green')
sns.histplot(data=df,x="temp_max",kde=True,ax=axs[0,1],color='red')
sns.histplot(data=df,x="temp_min",kde=True,ax=axs[1,0],color='skyblue')
sns.histplot(data=df,x="wind",kde=True,ax=axs[1,1],color='orange')
```

```
Out[85]: <AxesSubplot:xlabel='wind', ylabel='Count'>
```



```
In [87]: df.head()
```

```
Out[87]:
```

	precipitation	temp_max	temp_min	wind	weather
0	0.000000	12.8	5.0	2.167948	drizzle
2	0.894427	11.7	7.2	1.516575	rain
4	1.140175	8.9	2.8	2.469818	rain
5	1.581139	4.4	2.2	1.483240	rain
6	0.000000	7.2	2.8	1.516575	rain

```
In [88]: df.columns
```

```
Out[88]: Index(['precipitation', 'temp_max', 'temp_min', 'wind', 'weather'], dtype='object')
```

```
In [89]: X = df[['precipitation', 'temp_max', 'temp_min', 'wind']].values  
X[0:5]
```

```
Out[89]: array([[ 0.          , 12.8          ,  5.          ,  2.16794834],  
       [ 0.89442719, 11.7          ,  7.2          ,  1.51657509],  
       [ 1.14017543,  8.9          ,  2.8          ,  2.46981781],  
       [ 1.58113883,  4.4          ,  2.2          ,  1.4832397 ],  
       [ 0.          ,  7.2          ,  2.8          ,  1.51657509]])
```

```
In [91]: Y=df['weather'].values  
Y[0:5]
```

```
Out[91]: array(['drizzle', 'rain', 'rain', 'rain', 'rain'], dtype=object)
```

```
In [92]: X=preprocessing.StandardScaler().fit(X).transform(X.astype(float))  
X[0:5]
```

```
Out[92]: array([[ -0.60933199, -0.5794075 , -0.66479802,  1.34668256],  
       [  0.57361988, -0.72521075, -0.23932352, -0.49985515],  
       [  0.89864167, -1.09634631, -1.09027252,  2.20243364],  
       [  1.48185123, -1.69281417, -1.20631102, -0.59435558],  
       [ -0.60933199, -1.32167861, -1.09027252, -0.49985515]])
```

```
In [111... x_train,x_test,y_train,y_test=train_test_split(X,Y,test_size=0.1,random_state=2)  
print('Train set:',x_train.shape,y_train.shape)  
print('Test set:',x_test.shape,y_test.shape)
```

```
Train set: (1109, 4) (1109,)  
Test set: (124, 4) (124,)
```

```
In [112... knn=KNeighborsClassifier()  
knn.fit(x_train,y_train)  
print("KNN Accuracy:{:.2f}%".format(knn.score(x_test,y_test)*100))
```

```
KNN Accuracy:84.68%
```

```
In [96]: K=4  
neigh=KNeighborsClassifier(n_neighbors=K).fit(x_train,y_train)  
neigh
```

```
Out[96]: KNeighborsClassifier(n_neighbors=4)
```

```
In [97]: yhat=neigh.predict(x_test)  
yhat[0:5]
```

```
Out[97]: array(['sun', 'snow', 'sun', 'sun', 'snow'], dtype=object)
```

```
In [98]: print('Train set Accuracy:',metrics.accuracy_score(y_train,neigh.predict(x_train)))
```



```
print('Test set Accuracy:', metrics.accuracy_score(y_test, yhat))
```

Train set Accuracy: 0.842200180342651

Test set Accuracy: 0.8225806451612904

```
In [113... ks=10
mean_acc = np.zeros((ks-1))
std_acc = np.zeros((ks-1))

for n in range(1, ks):
    neigh=KNeighborsClassifier(n_neighbors=n).fit(x_train, y_train)
    yhat=neigh.predict(x_test)
    mean_acc[n-1]=metrics.accuracy_score(y_test, yhat)
    std_acc[n-1]=np.std(yhat==y_test)/np.sqrt(yhat.shape[0])

mean_acc
```

```
Out[113]: array([0.7983871 , 0.71774194, 0.81451613, 0.82258065, 0.84677419,
                0.86290323, 0.86290323, 0.83064516, 0.83064516])
```

```
In [106... from sklearn.preprocessing import LabelEncoder
```

```
In [107... lc=LabelEncoder()
df["weather"]=lc.fit_transform(df["weather"])
```

```
In [108... df.head()
```

```
Out[108]:
```

	precipitation	temp_max	temp_min	wind	weather
0	0.000000	12.8	5.0	2.167948	0
2	0.894427	11.7	7.2	1.516575	2
4	1.140175	8.9	2.8	2.469818	2
5	1.581139	4.4	2.2	1.483240	2
6	0.000000	7.2	2.8	1.516575	2

```
In [115... x=((df.loc[:, df.columns!="weather"]).astype(int)).values[:,0:]
y=df["weather"].values
```

```
In [110... df.weather.unique()
```

```
Out[110]: array([0, 2, 4, 3, 1])
```

```
In [103... !pip install xgboost
```

Collecting xgboost

Downloading xgboost-1.7.2-py3-none-win_amd64.whl (89.1 MB)

Requirement already satisfied: numpy in c:\users\ahmed\anaconda3\lib\site-packages (from xgboost) (1.21.5)

Requirement already satisfied: scipy in c:\users\ahmed\anaconda3\lib\site-packages (from xgboost) (1.7.3)

Installing collected packages: xgboost

Successfully installed xgboost-1.7.2

```
In [104... from xgboost import XGBClassifier
```

```
In [127... le = LabelEncoder()  
y_train = le.fit_transform(y_train)
```

```
In [124... import warnings  
warnings.filterwarnings('ignore')  
xgb=XGBClassifier()  
xgb.fit(x_train,y_train)  
print("XGB Accuracy:{:.2f}%".format(xgb.score(x_test,y_test)*100))
```

XGB Accuracy:0.00%

```
In [128... input=[[1.140175,8.9,2.8,2.469818]]  
ot=xgb.predict(input)  
print("The weather is:")  
if(ot==0):  
    print("Drizzle")  
elif(ot==1):  
    print("Fog")  
elif(ot==2):  
    print("Rain")  
elif(ot==3):  
    print("snow")  
else:  
    print("Sun")
```

The weather is:

Rain

```
In [ ]:
```